

Flame: A distributed system for intelligent workload

Klaus Ma (@k82cn, Nvidia)



Content

目录

01 Why Flame?

02 Use Cases

03 Architecture

04 Roadmap


AI

Why Flame?



- Currently, more and more frameworks are introduced for specific domain, e.g., BigData, AI; but the meta operations are similar, e.g., matrix, gradient, map/reduce.
- Meanwhile, no meta framework for specific distributed system use cases, e.g., Monte Carlo, Crawler, Encoder.
- Flame is a distributed engine for the frameworks in domain; the Frame focus on the meta-API, performance, throughput, scheduling and high availability.
- Exploring a new way of data sharing/exchanging

Example: Pi by Monte Carlo




Pi Client:

1. Create a callback for each task
2. Create tasks based on input
3. Waiting for all tasks completion
4. Print the estimation of Pi

Pi Service:

1. Generate random points by input
2. Print how many points are in the circle

Example: Matrix Multiplication



Matrix Client:

1. Create a callback for each tasks
2. Create tasks based on input
3. Waiting for all tasks completion
4. Print all items of the matrix

Matrix Service:

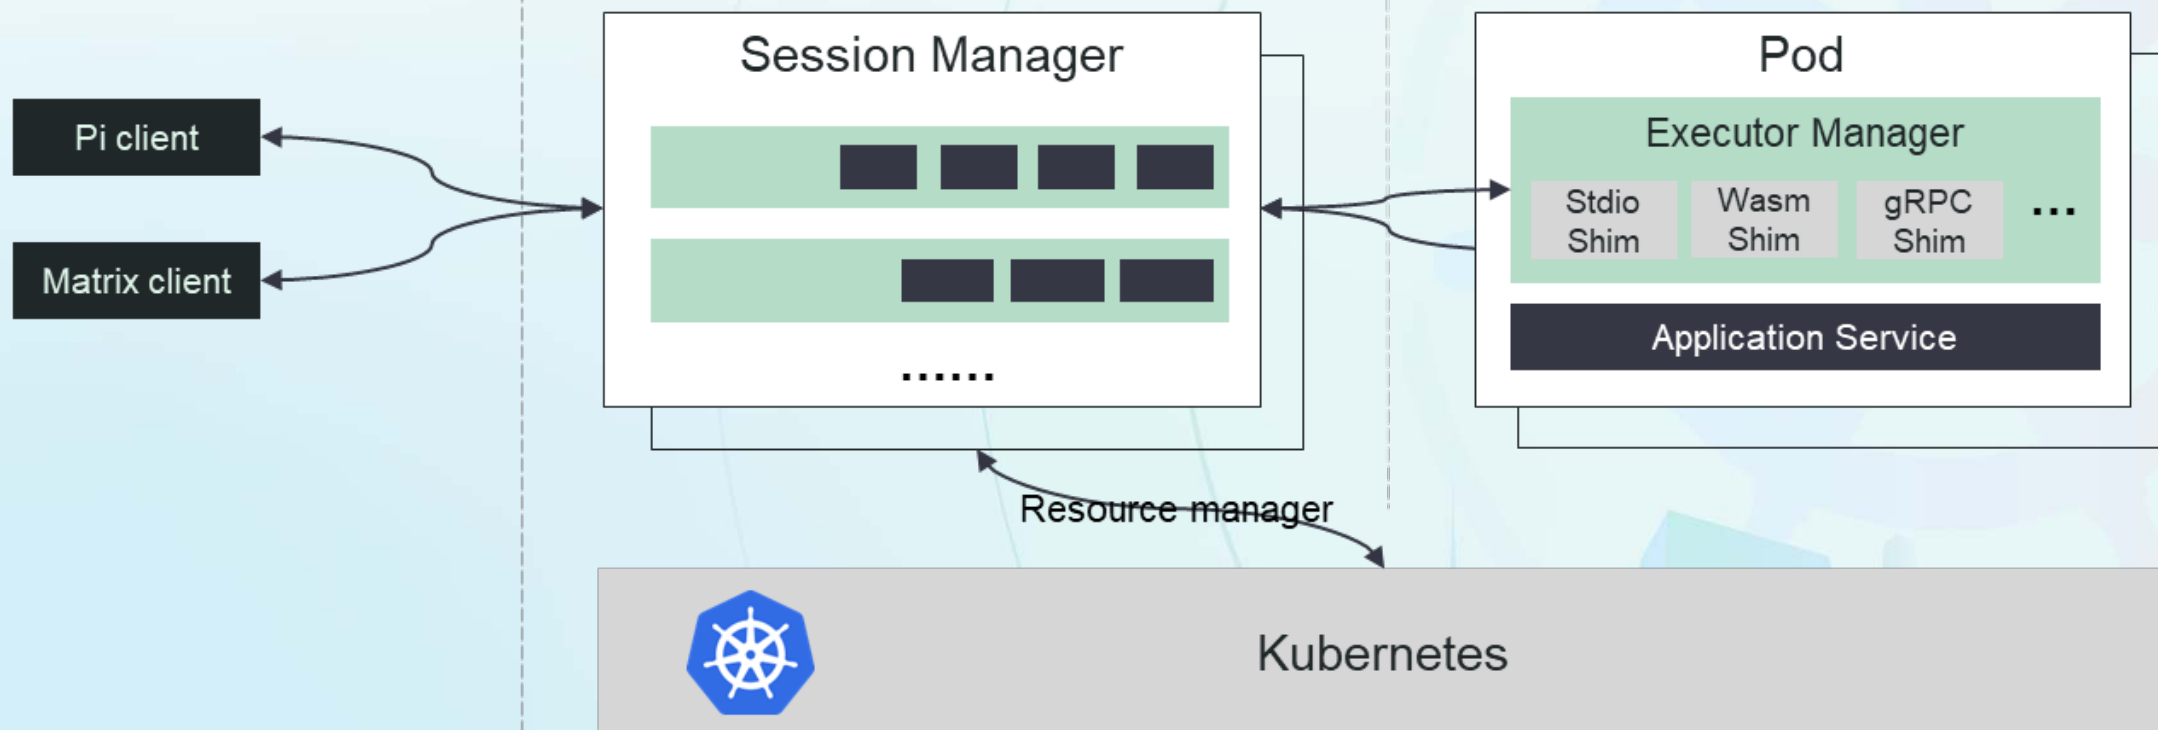
1. Calculate C_{ij} for as the result

Overall Architecture

Application Client
(Flame API, e.g. Python, Rust)

Flame Session Manager
(e.g. faire-share, DAS, priority)

Flame Executor Manager
(e.g. gRPC, Wasm, process)



Flame SDK (core)

Flame Client

`flame.connect()`

`flame.create_session()`

`flame.run_task(callback)`

`flame.close_session()`

`flame.disconnect()`

Flame

Flame Service

`flame.on_session_enter()`

`flame.on_task_invoke()`

`flame.on_session_leave()`

Flame Python SDK (on-going)


Flame Client

`flame.init()`

`@flame.service`

`flame.future`

Python bytecode



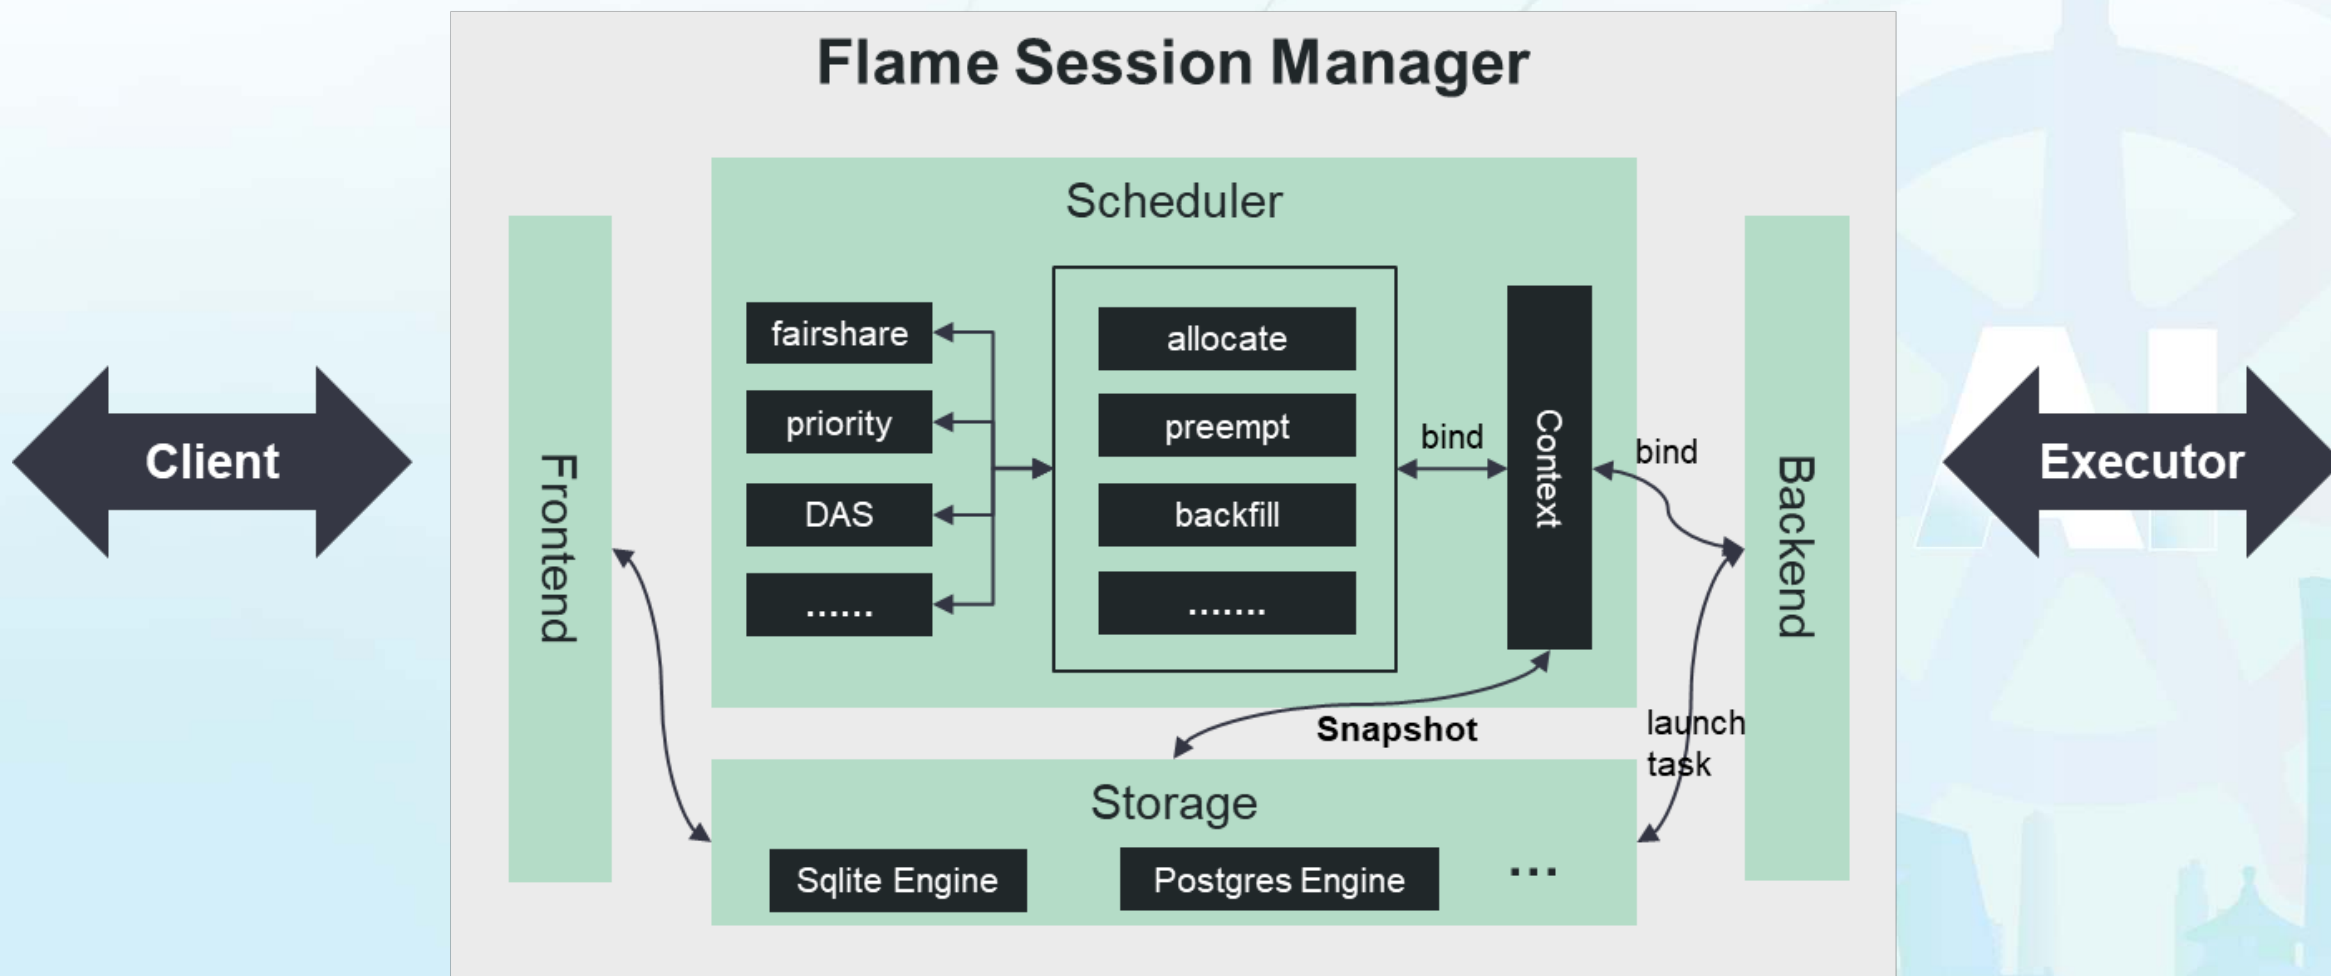
Flame Python by gRPC

`flame.on_session_enter()`

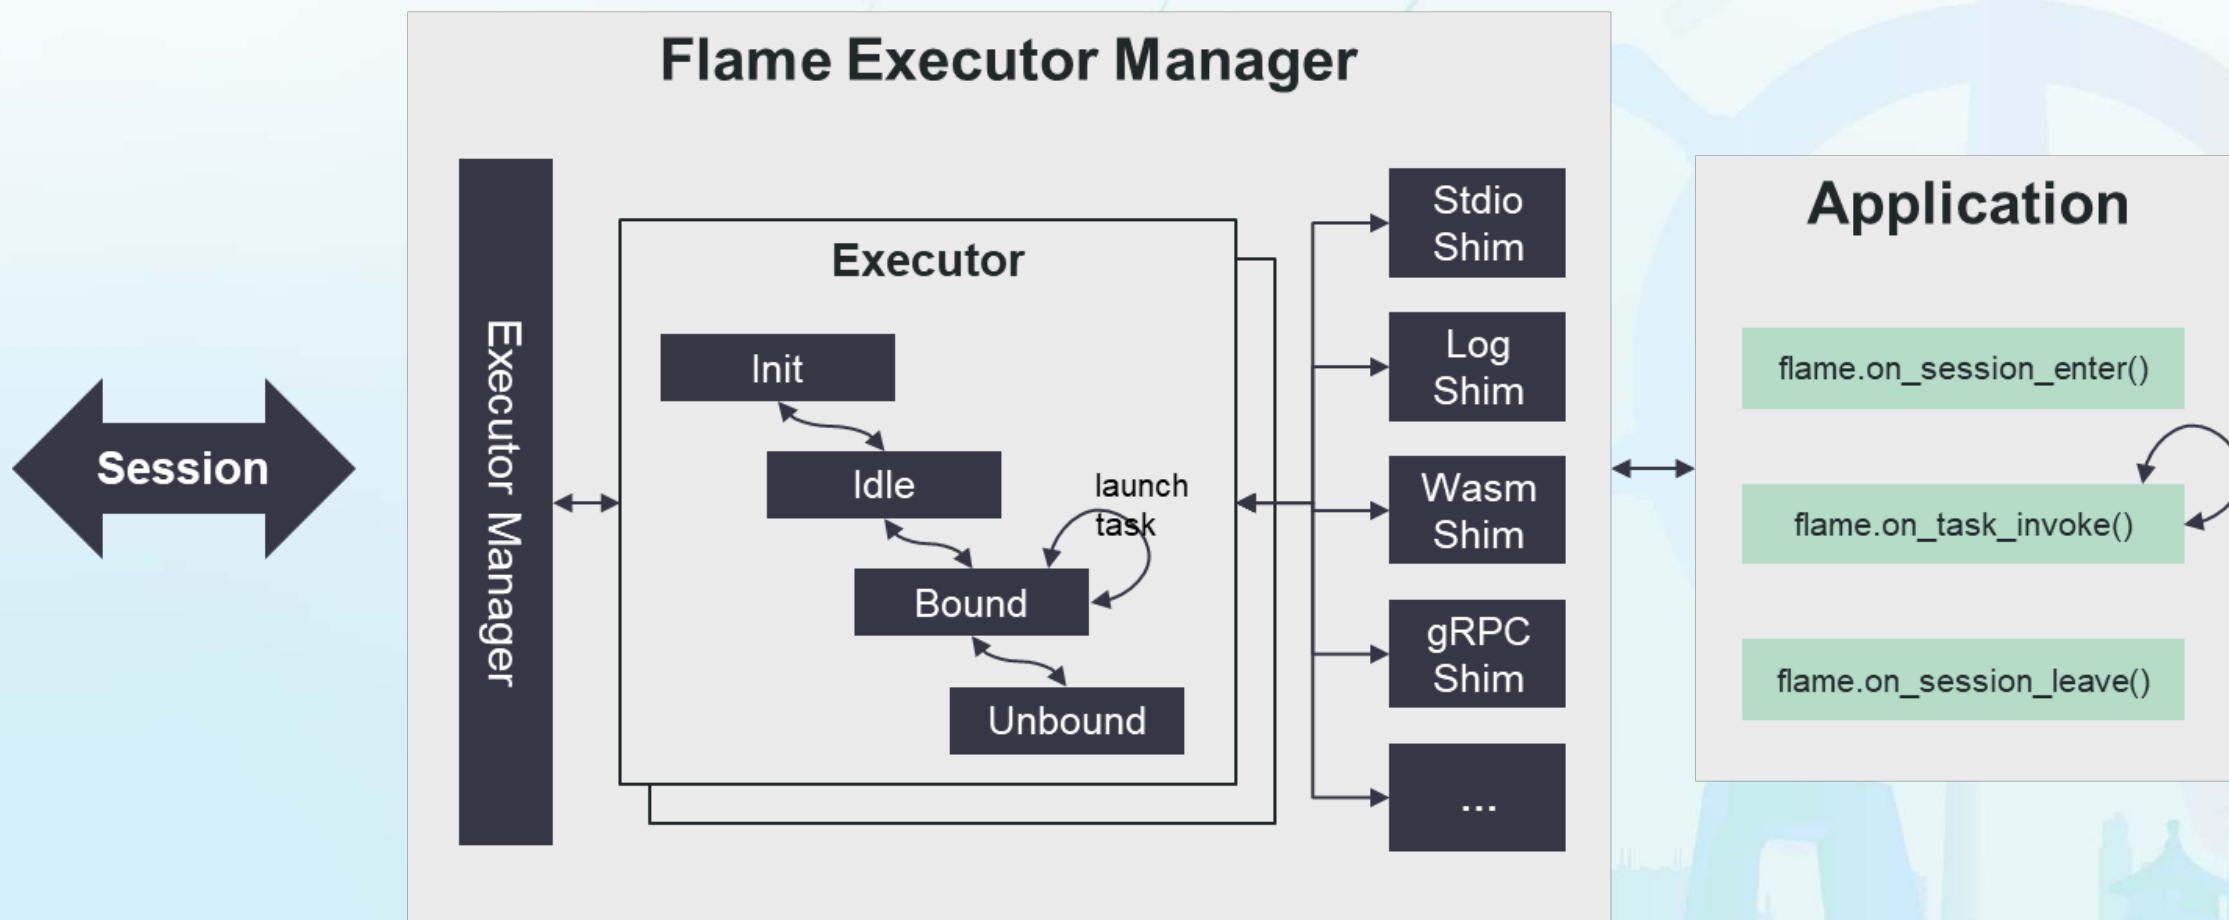
`flame.on_task_invoke()`

`flame.on_session_leave()`

Session Manager



Executor Manager



Roadmap

- **gRPC shim for all language, e.g. Python, R, C++ (Done)**
- Distributed storage/cache for **data sharing/exchanging**
- PyTorch, Tensorflow examples, e.g. distributed training
- TLS/mTLS enhancement for all connection
- More scheduling policy, e.g. priority, minService/maxService, data aware scheduling
- Resource manager integration, e.g. Kubernetes
- Misc, e.g. CLI, matrices, dashboard
- flame-operator to simplify operations
- Documentation and tutorial

References



- [flame-sh/flame: A distributed system for intelligent workload \(github.com\)](#)
- [Monte Carlo method – Wikipedia](#)
- [Matrix multiplication - Wikipedia](#)

AI

Thanks.

