

# Exploring the Next Generation of Secure Containers: gVisor and Kata Fusion

**Xuewei Niu** Software Engineer, Ant Group

**Hang Su** Software Engineer, Ant Group

**Tiwei Bie\*** Staff Engineer, Ant Group

# Content

**01** Secure Containers Overview

**02** Inside Existing Solutions

**03** User-mode Linux

**04** NanoPVM

AI



# Part 01

## Secure Containers Overview

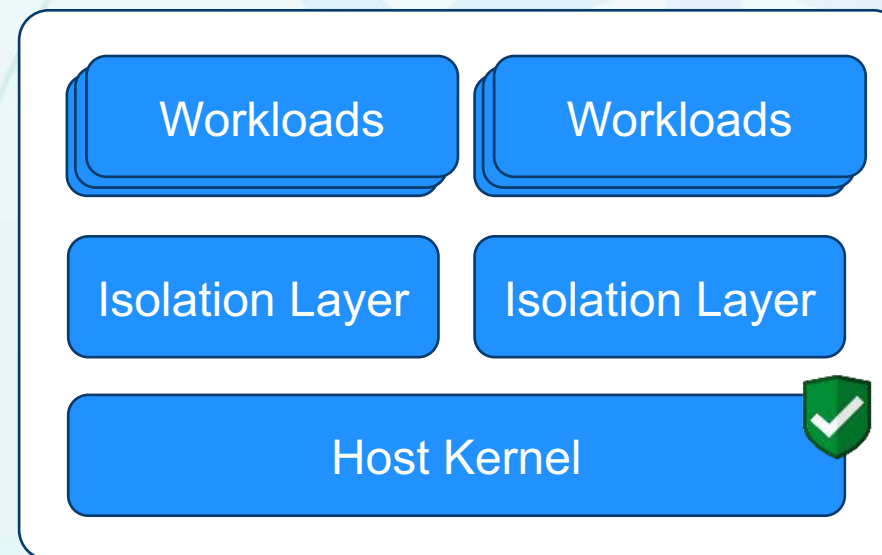
Why secure containers?



# Secure Containers Overview

Secure containers are more than security

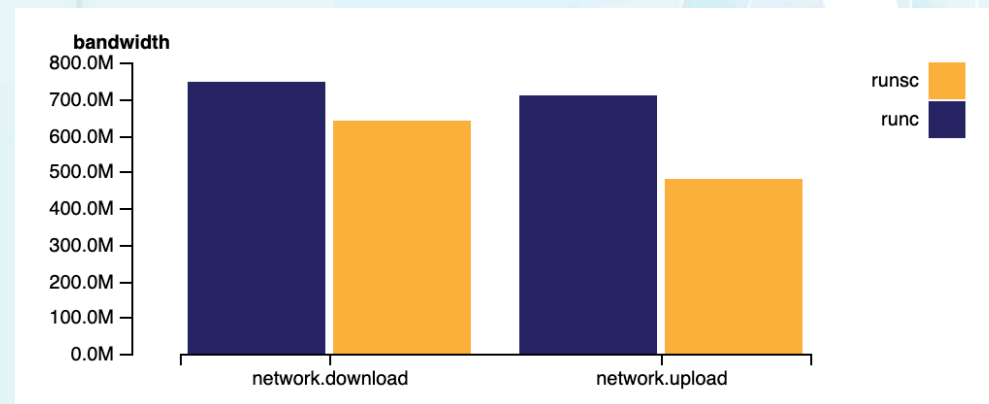
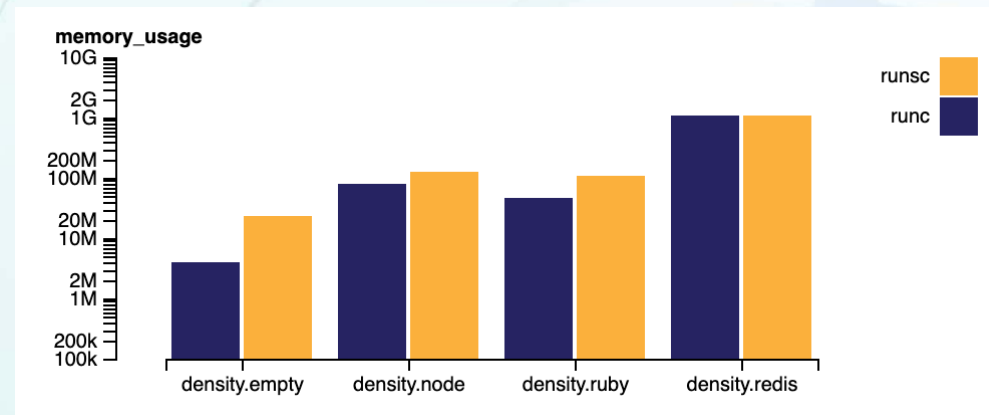
- **Security isolation** prevents sensitive instruction escapes
  - Executing untrusted code
  - Multi-tenancy
- **Performance isolation** prevents scheduling, networking, I/O interference
  - Online-offline hybrid deployment
- **Fault isolation** prevents shared kernel crashes/faults



# Secure Containers Overview

At the mean time...

- Runtime overheads cause increased latency, reduced throughput or density
- Resource footprints from additional components, e.g. guest kernel/Sentry



\*Data from gVisor Documentation  
[https://gvisor.dev/docs/architecture\\_guide/performance/](https://gvisor.dev/docs/architecture_guide/performance/)

# Part 02

## Inside Existing Solutions

Brief Introduction to gVisor and Kata Containers





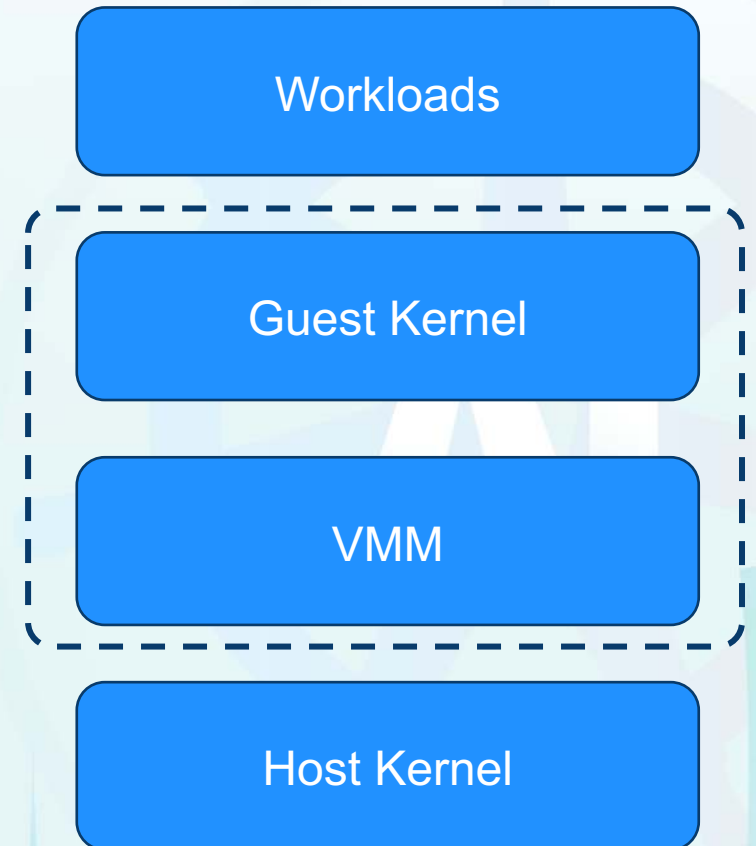
# Inside Existing Solutions

## Secure Containers Landscape



# Inside Existing Solutions: Kata Containers


- The speed of containers, the security of VMs
- Kata provides a **virtual machine**: a typical model with a clear layered structure
  - VMM provides virtual devices
  - Guest kernel provides dedicated runtime environment
- Most of the work is focused on reducing overheads
  - MicroVMs: Firecracker (for FaaS), Cloud-hypervisor (for general tasks), etc.
  - Disabled unnecessary kernel configs






# Inside Existing Solutions: gVisor

- gVisor is an independent kernel running in user-space, exposing a *Linux-like* interface.
- gVisor provides a **virtual kernel** with merged VMM and guest kernel layers. The model eliminates unnecessary virtualization assumptions.
- Supported platforms: KVM, systrap and ptrace
- Sentry  $\approx$  kernel: When the application makes a system call, the platform redirects the call to the Sentry, which will do the necessary work to service it.



# Inside Existing Solutions

 **RUNC**



?

So... what is the next?

# Part 03

## User-Mode Linux

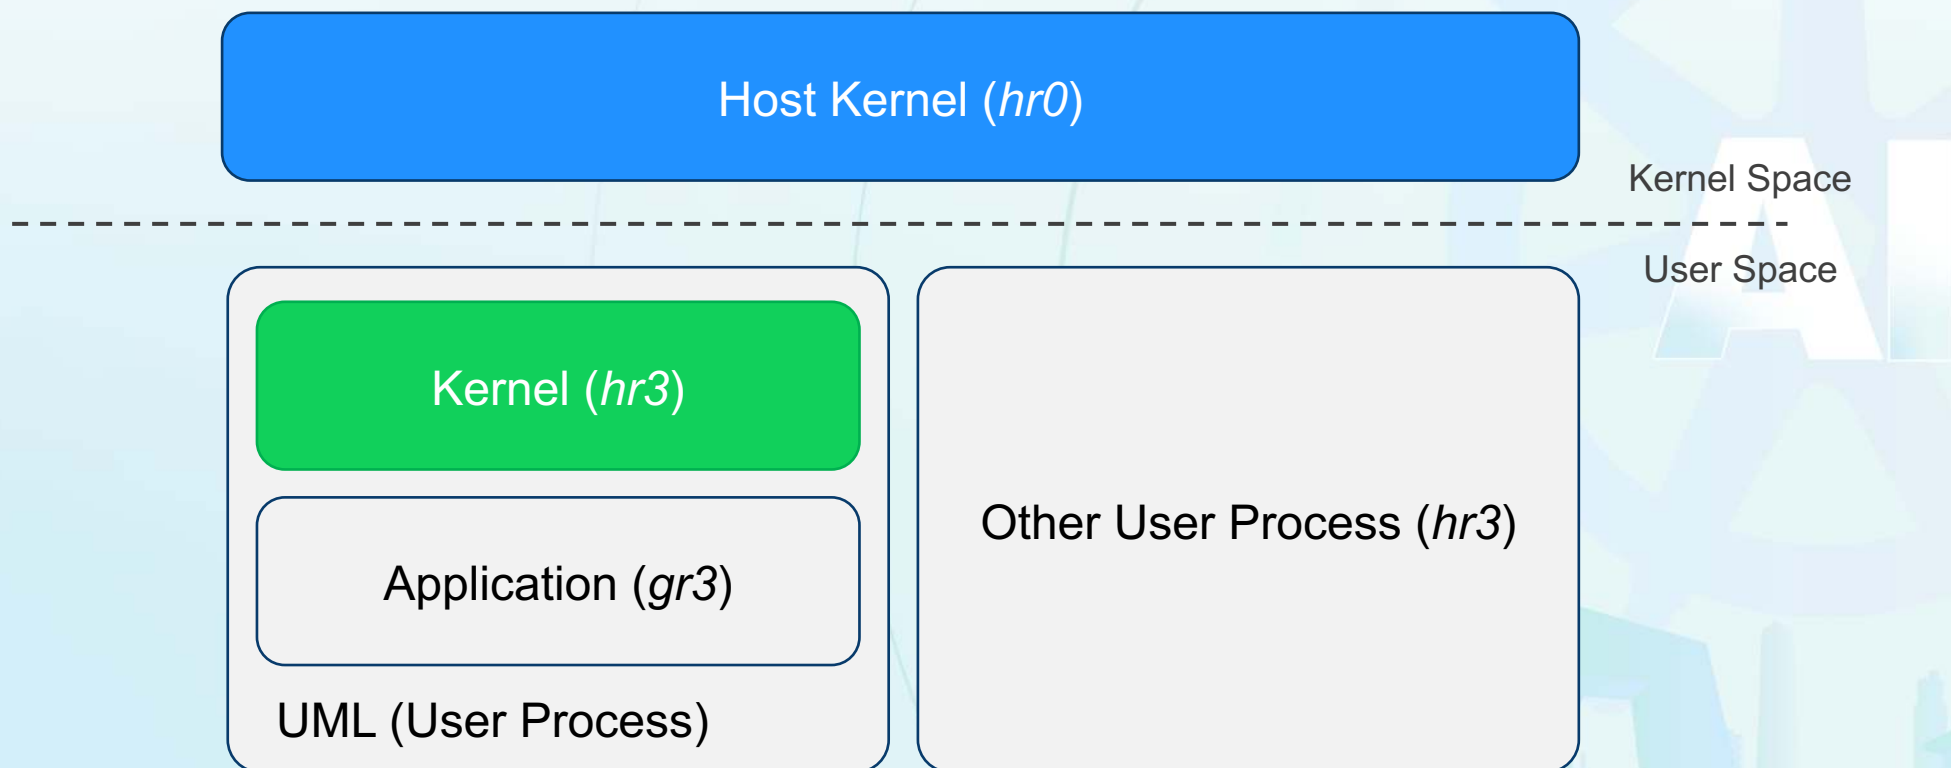
The Next Generation of Secure Containers



# User-Mode Linux


UML was first released in 1991 and has been merged into the Linux mainline.

- Build with `make ARCH=um`



# User-Mode Linux

UML is more like gVisor from perspective of architecture



# User-Mode Linux

Compared with gVisor

- Running applications without modification or adaptation
- Time-tested kernel (also known as implementation costs)
- Costs of Golang: runtime, garbage collection, etc.
- Free to use any customized kernels

Compared with Kata Containers

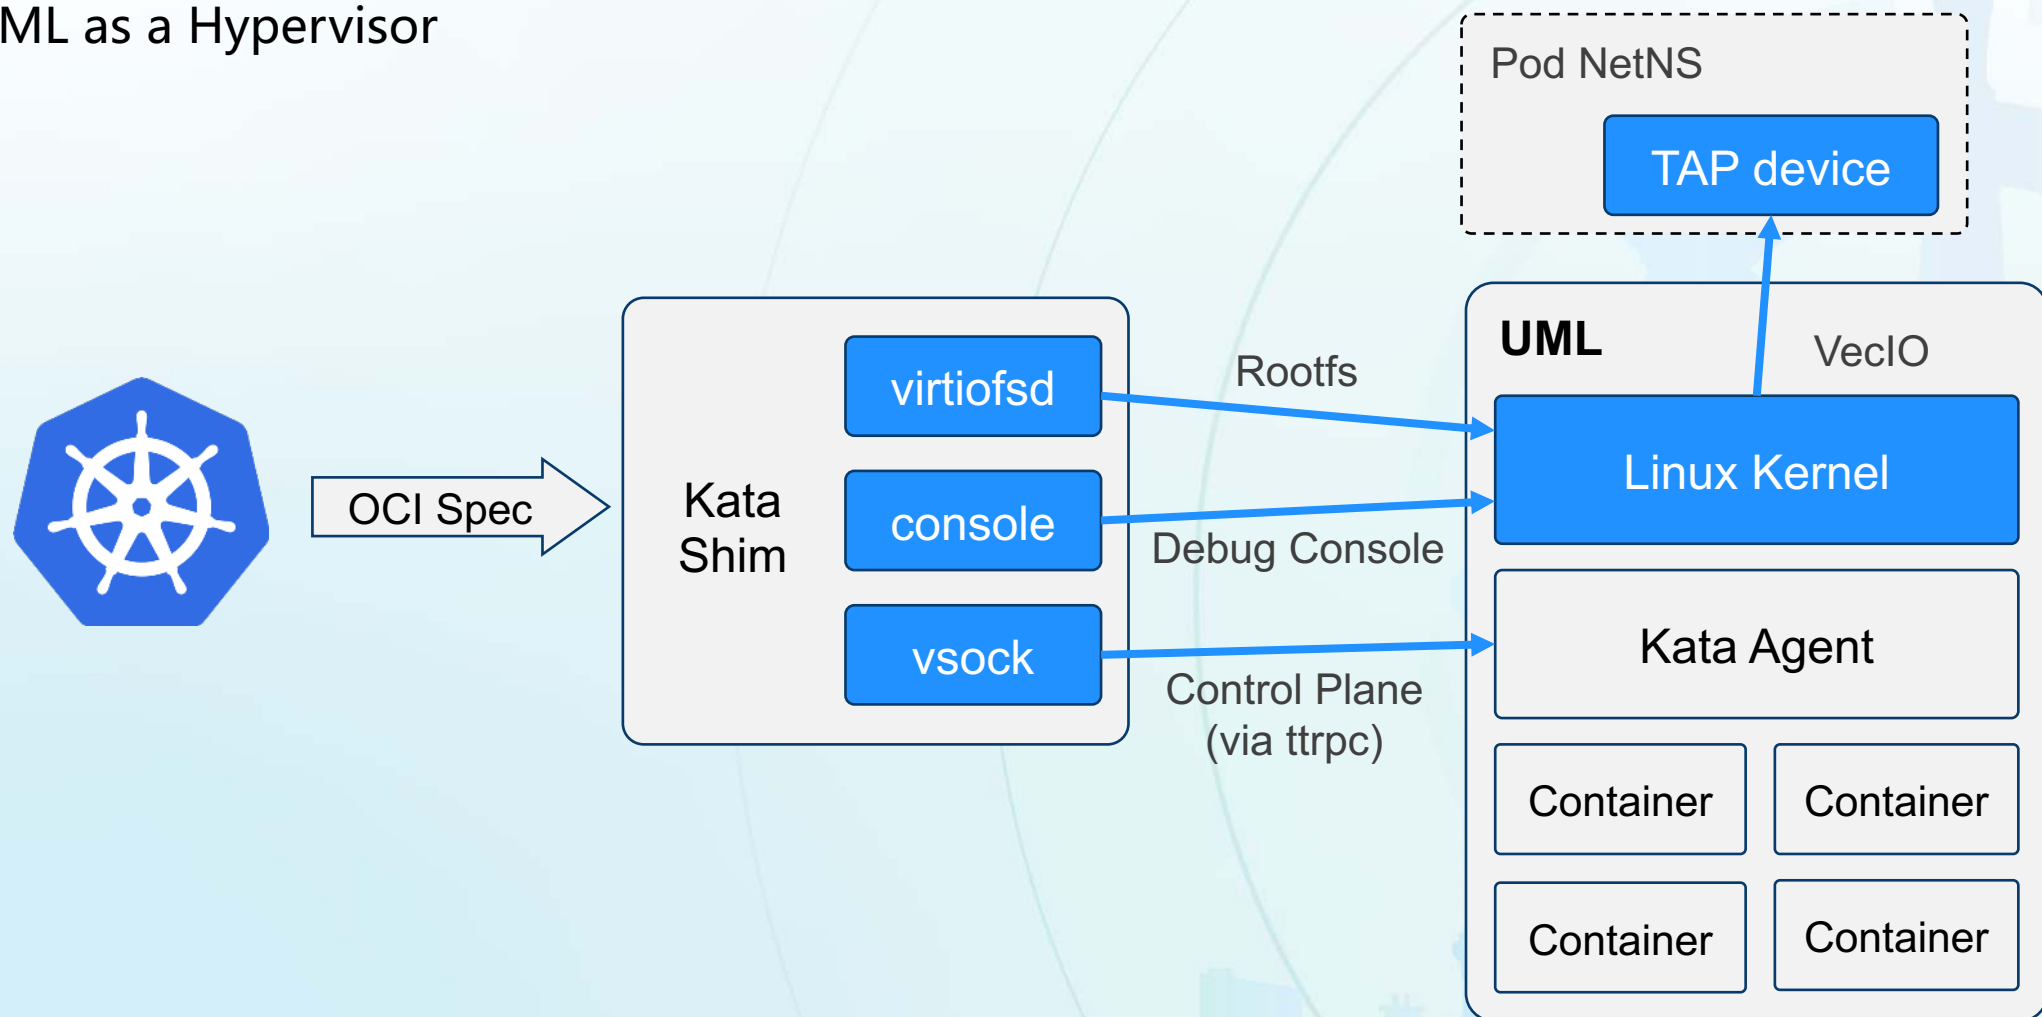
- With **NanoPVM**, running in environments that do not support nested virtualization
- Simpler virtualization model
- vGPU support (nvproxy)

AI



# User-Mode Linux

## UML as a Hypervisor



# Part 04

## NanoPVM

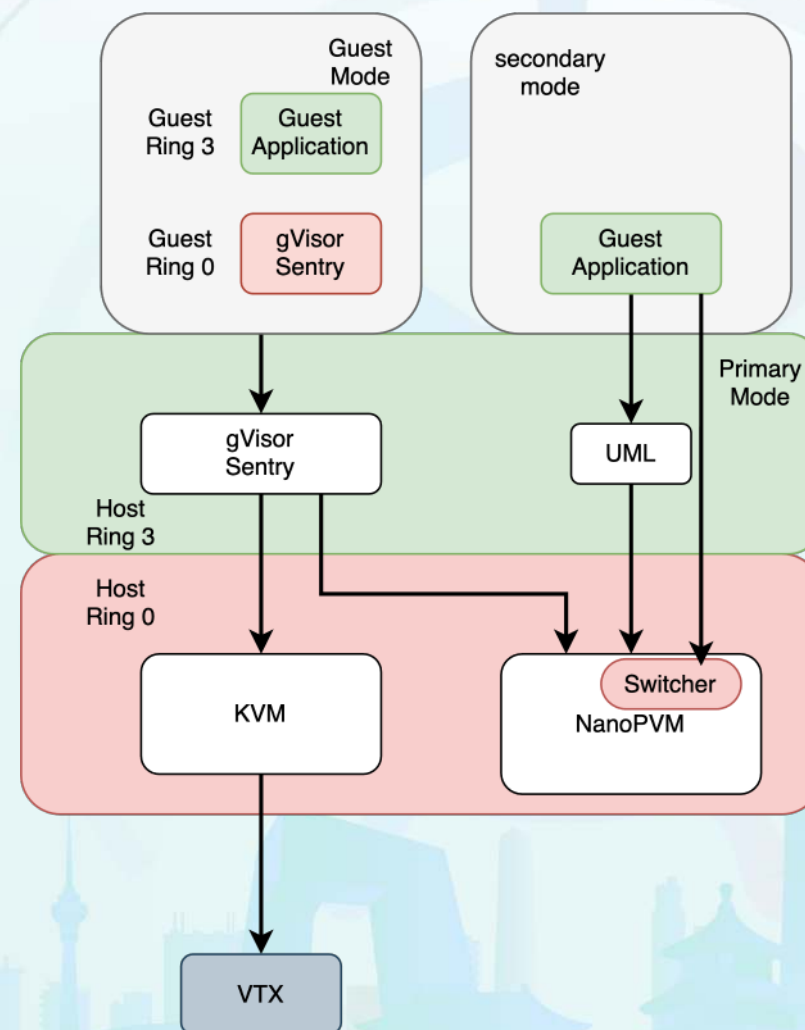
Enables User Mode Kernel to Run in Any Environments

AI



# What is NanoPVM ?

- Same architectural level as KVM ,  
Currently exists as an out-of-tree Linux kernel module
- **4 Core points:**
  - **New virtualization model paradigm**
  - Memory isolation utilizes **Direct Memory Management**
  - **Executable Context Management**, No instruction emulation
  - Minimalist Interface



# New virtualization model paradigm

Popek and Goldberg virtualization requirements:

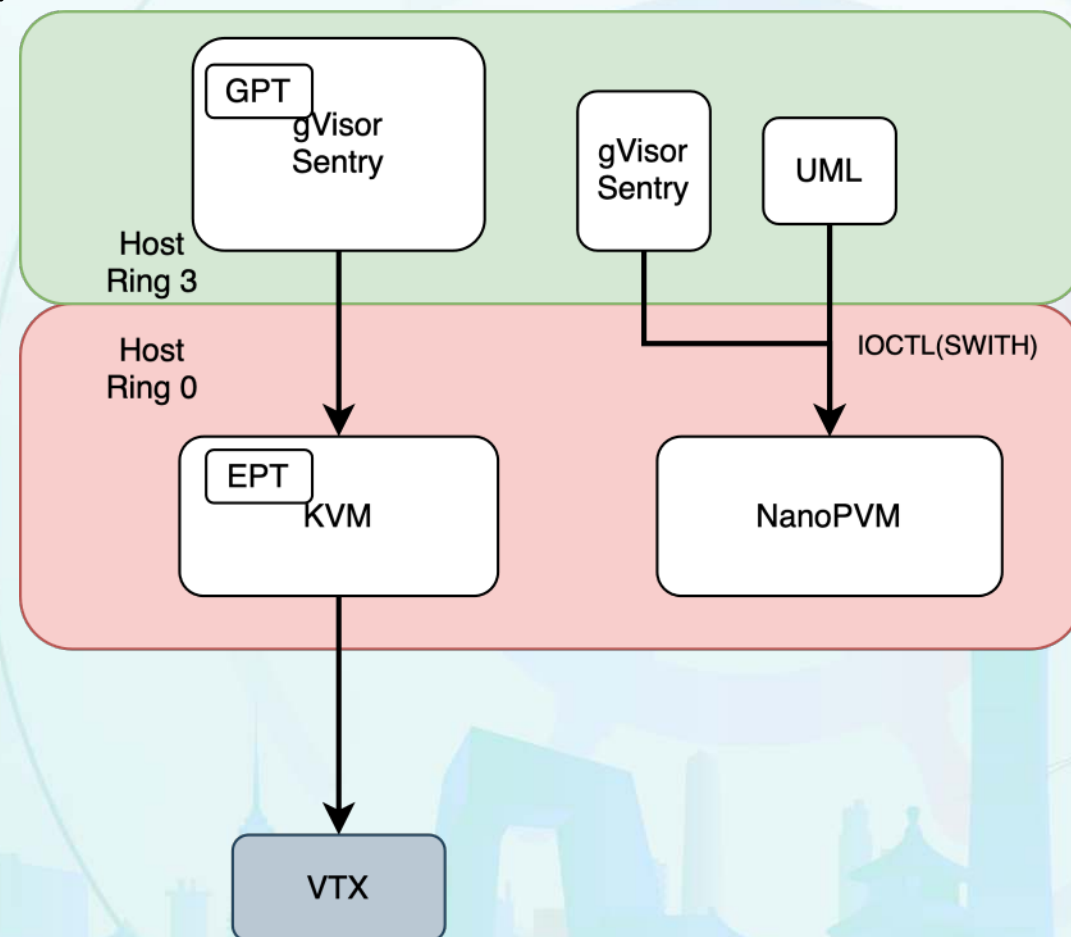
- Equivalence / Fidelity
- Resource control / Safety
- Efficiency / Performance
- Provide 2 Virtualization ability: **Context Execution, Address Space Management**
- Implementation with 3 Control Interfaces : **Map, Unmap, Switch**

# OS Interface Complete

OS Function	Implement method	Interface
Create Process	Allocate address space and setup initial process status	map + switch
Process Exit	Release all memory mappings and resource	unmap
Memory Manage	Map/Unmap mapping on demand	map + unmap
Syscall Handle	Trap and handle after VM-Exit	switch
Exception Handle	Trap and handle after VM-Exit	switch
Device IO	Map device memory or port to host	map + switch

# Secondary Mode

- Direct Memory Management via Map and Unmap.
- Intercept Syscall and Exception to User Mode Kernel.





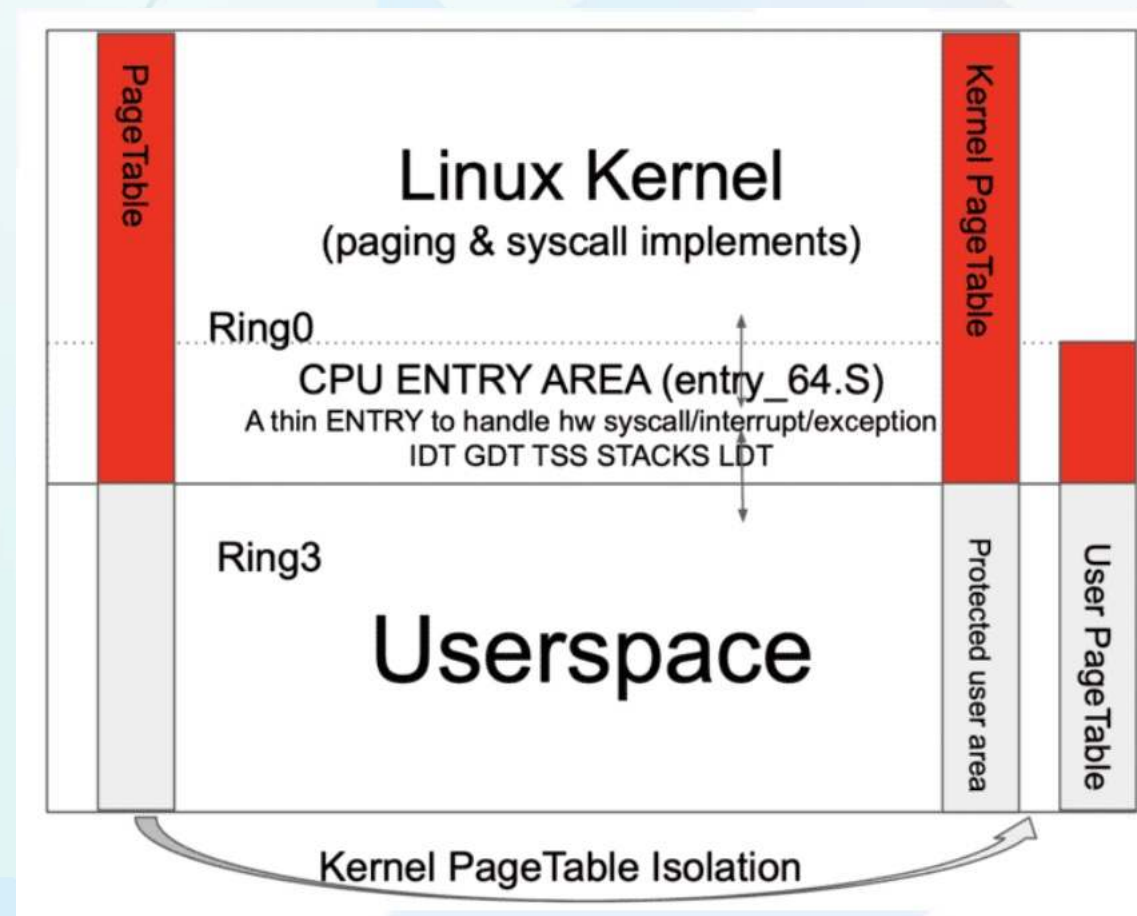
# Interface and Usage example

```
/* Create compute instance: */  
  
instance_fd = open("/dev/nanopvm");  
  
as_fd = ioctl(instance_fd, NANOPVM_AS_CREATE);  
  
ioctl(asfd, NANOPVM_MAP, *mem_range);  
  
  
/* Virtualization loop: */  
  
do {  
    ioctl(asfd, NANOPVM_SWITCH, *context_regs);  
    ...  
    handle_vm_exit();  
} while(1);  
  
  
/* Process exit: */  
  
ioctl(asfd, NANOPVM_UNMAP, *mem_range);
```

AI

# Syscall/Exception Intercept

- Intercept without Ptrace, but **hook the trampoline entry**.
- Provide “Executable” context, and 4 capabilities:
  - enter vm
  - force kick out (bounce)
  - exception exit
  - syscall exit
- With the help of **Switcher**



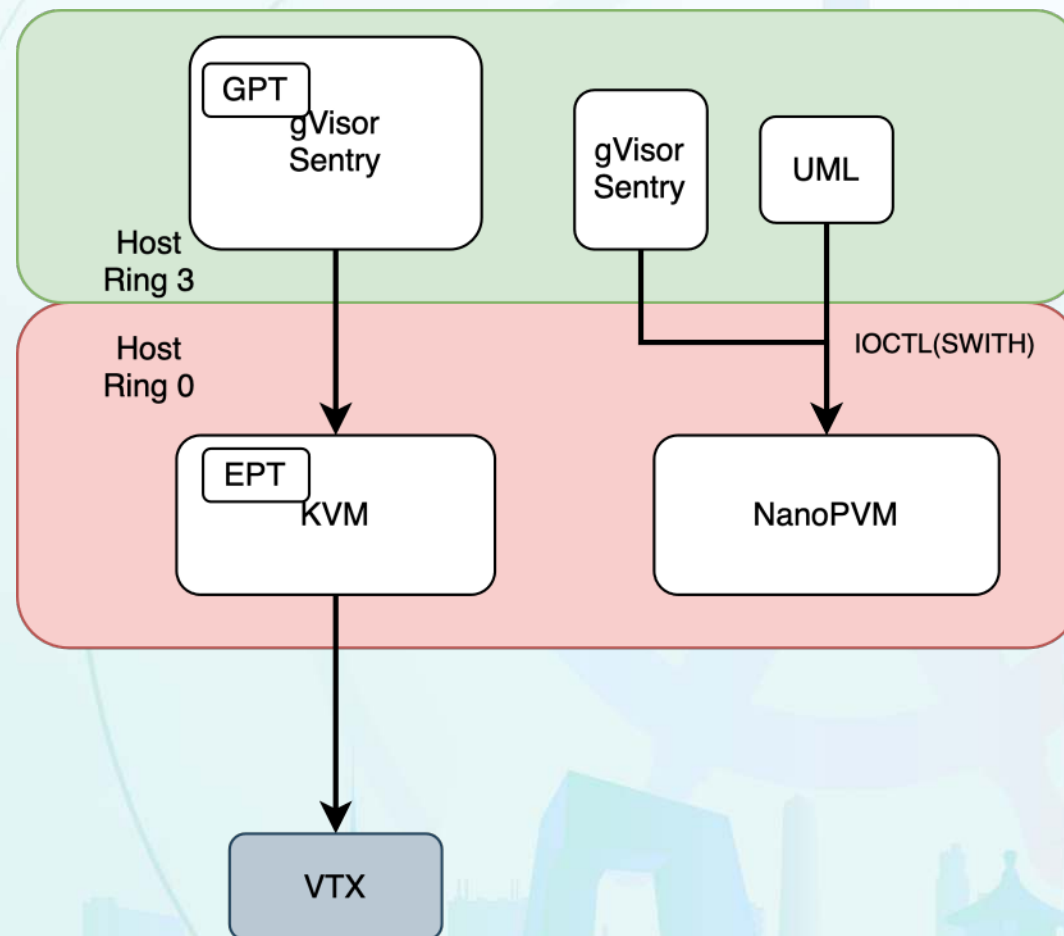
# Secondary Mode Switch

- Called with context\_switch by UML on purpose
- Call with Guest\_Context\_Regs {

// fields similar to ptrace\_regs

.....

}



# Memory Management

- Called from mmap, brk, handle\_user\_fault and so on.
- Call with Guest\_Memory\_Range {

void \*gva;

void \*hva;

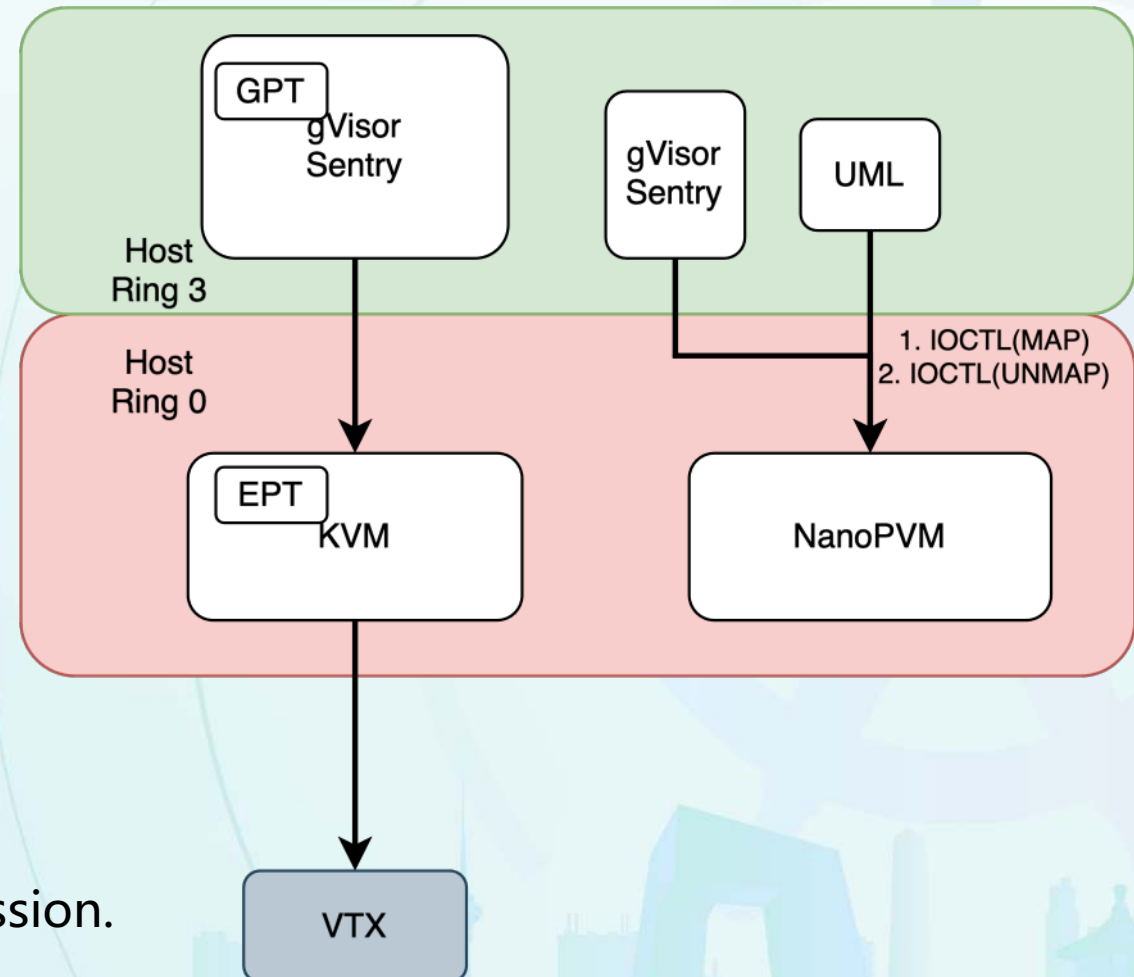
uint len;

uint prot;

.....

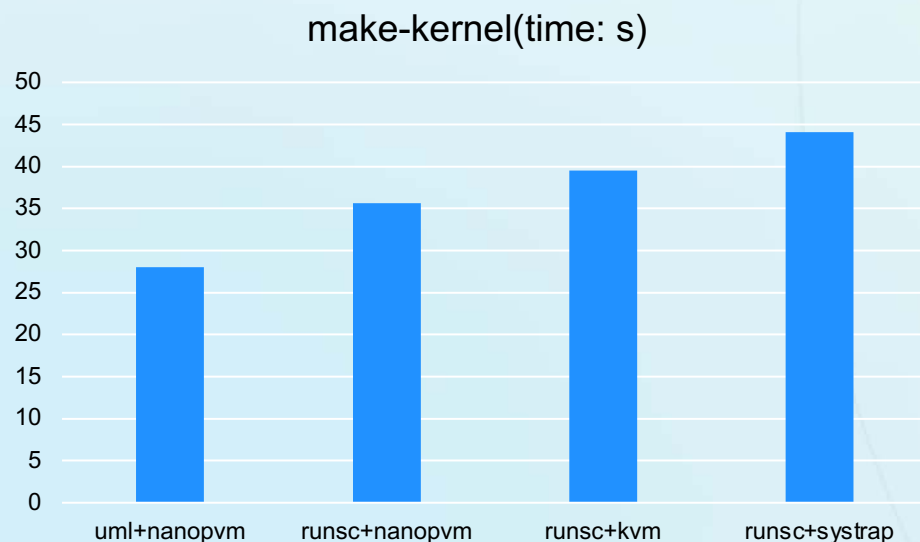
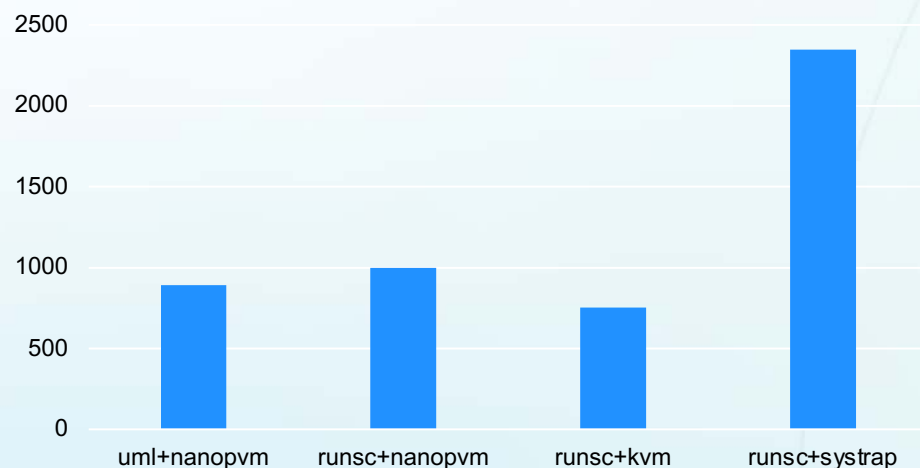
}

- **No conception of GPT** at all.
- All Guest address mapped with User permission.
- Memory region Management with rbTree.




# Performance

Compared with UML+NanoPVM, gVisor+NanoPVM, gVisor+KVM, gVisor-systrap  
getpid(time:ns)



## Redis performance



# Future

- The basement of Next Security Container.
- Alternative base tech for other kind of Sandbox, e.g. WebAssembly Runtime.
- (Maybe) Remove the GPT from UML.

AI



# Thanks and QA